

Learning to Distrust Misleading Text: Preference Optimization for Reducing Text Bias in Vision-Language Models

Anonymous CVPR submission

Paper ID *****

Abstract

001 Vision-language models (VLMs) often treat textual context
002 as a shortcut, even when that text conflicts with visual evi-
003 dence. This “blind faith in text” is especially problematic
004 in document and visual question answering, where auxil-
005 iary text may be incomplete, stale, adversarial, or automa-
006 tically retrieved. We study whether preference optimization
007 can reduce this failure mode without sacrificing normal vi-
008 sual question answering performance. The goal is not to
009 teach the model to ignore all text: auxiliary text is often
010 useful when it agrees with the image, and the model usually
011 does not know in advance whether the text is reliable. Start-
012 ing from corrupted-text evaluations, we mine model-specific
013 text-following errors: cases where the model is wrong and
014 its wrong answer is copied from the corrupted auxiliary
015 text. We then train VLMs with Direct Preference Optimiza-
016 tion (DPO), preferring the image-grounded answer over
017 the model’s own text-following error. On DocVQA with
018 Qwen3-VL-8B, text-following span DPO substantially im-
019 proves corrupted-text robustness and reduces the fraction
020 of remaining errors copied from misleading text. A multi-
021 condition error-mined extension further strengthens this ef-
022 fect. By contrast, GRPO-style sampled-answer optimiza-
023 tion and auxiliary SFT preservation do not improve over
024 DPO. Experiments on VQAv2 and transfer to Qwen2-VL
025 and LLaVA-Next show that the method works best when the
026 base VLM already has sufficient visual and OCR compe-
027 tence. These results suggest that high-precision preference
028 pairs are an effective and practical mechanism for correct-
029 ing text bias in VLMs, while also revealing clear boundary
030 conditions for the approach.

031 1. Introduction

032 Vision-language models (VLMs) are increasingly used in
033 settings where images are accompanied by textual context:
034 document images with OCR, screenshots with metadata, re-
035 trieved descriptions, captions, or user-provided hints. Such

text can be helpful when it is correct, but it can also be-
come a shortcut. When textual context conflicts with the
image, a VLM may answer from the text rather than from
visual evidence. Deng et al. [3] describe this phenomenon
as “blind faith in text”, showing that several VLMs dispro-
portionately trust textual information under conflict.

This failure is not merely a benchmark artifact. In prac-
tical systems, auxiliary text is often produced by retrieval,
OCR, captioning, or prior model predictions. All of these
sources can be noisy. If a VLM over-trusts such text, the
model can produce confident but visually ungrounded an-
swers. For document VQA, this is particularly important:
the model must decide whether the image or the accompa-
nying text should be trusted when they disagree.

A natural objection is that the model could simply ignore
the auxiliary text whenever corruption is possible. This is
not a satisfactory solution. First, in real use the system is
not told whether the text is corrupted, matched, or irrele-
vant. Second, auxiliary text is often genuinely useful: when
it matches the image, ignoring it can remove information
that helps the model answer. Third, blind text removal does
not teach conflict resolution; it only changes the input dis-
tribution. The desired behavior is conditional trust: use text
when it is supported by the image, but reject it when it con-
tradicts visual evidence.

We ask a targeted question: *Can we reduce text-
following failures by directly training a VLM to prefer
image-grounded answers over its own corrupted-text er-
rors?* Our key observation is that the corrupted-text setting
naturally produces preference pairs. If a model answers in-
correctly and its prediction appears in the corrupted auxil-
iary text, then the model has likely followed the mislead-
ing text. The ground-truth answer can be used as the pre-
ferred response, and the model’s text-following prediction
can be used as the dispreferred response. This gives a high-
precision preference signal without manual labeling.

We instantiate this idea with Direct Preference Opti-
mization (DPO) [11]. Unlike online reinforcement learning
methods such as PPO or GRPO [12], DPO directly opti-
mizes pairwise preferences. This is well aligned with our

Mitigating Blind Faith in Text in Vision-Language Models via Direct Preference Optimization




Figure 1. Overview of our mitigation. We mine cases where the base VLM follows corrupted auxiliary text, then train a DPO preference that ranks the image-grounded answer above the model’s own text-following error.

setting: the contrastive pair is already available as {image-grounded answer, text-following wrong answer}. Prior work already showed that supervised fine-tuning (SFT) with text augmentation can mitigate blind faith in text [3]. Our goal is not to deny that result, but to ask whether a more targeted preference signal can address the same failure mode. Instead of training on a broad mixture of text-only, original VQA, match, corruption, and irrelevance samples, we mine the specific cases where the model follows corrupted text and directly contrast that answer with the ground truth. We compare DPO against GRPO-style sampled-answer optimization, auxiliary SFT preservation, model transfer, and dataset transfer.

Figure 1 summarizes our mitigation. We treat Words or Vision as evidence that the conflict exists, and focus on which training signal most directly changes the behavior. Removing text changes the input rather than teaching conflict resolution; generic SFT supplies correct answers without explicitly identifying the misleading answer; on-line RL depends on candidate sampling and reward shaping. We therefore use preference pairs derived from actual text-following failures, where the chosen and rejected answers instantiate the desired behavioral reversal.

Our study makes three empirical findings. First, on DocVQA with Qwen3-VL-8B, DPO sharply improves corrupted-text robustness while reducing the rate of text-copying among remaining errors. Second, broadening DPO from corrupted-only text-following errors to multi-

condition error-mined pairs yields the strongest result, suggesting that condition-specific error mining can improve stability. Third, the method has an important boundary condition: it works best when the base VLM already has enough visual/OCR competence. Qwen-family models benefit substantially, while LLaVA-Next models under our prompt and template show weak DocVQA base performance and limited benefit.

We summarize our contributions as follows:

- We define an automatic text-following error mining procedure for corrupted-text VQA, requiring both incorrectness and overlap between the wrong prediction and corrupted auxiliary text.
- We show that DPO on these mined pairs substantially reduces text bias on DocVQA and transfers directionally to VQAv2 and Qwen2-VL.
- We position DPO relative to the SFT mitigation studied by Deng et al. [3], and compare it with GRPO, auxiliary SFT preservation, multi-condition error mining, and model transfer.

2. Related Work

Text bias in VLMs. VLMs are trained to fuse visual and textual inputs, but fusion can become asymmetric when textual context is easier to exploit than visual evidence. The most directly related work is Words or Vision [3], which introduces corrupted textual variations and shows that VLMs

may trust text over vision when the two conflict. Our work follows this diagnostic setting but focuses on mitigation through preference optimization. Rather than only measuring the degradation under corrupted text, we use the model’s own corrupted-text failures as training signal. Deng et al. [3] also study SFT with text augmentation and show that increasing SFT data reduces reliance on corrupted or irrelevant text while preserving the ability to use matched text. They further note that text-only data helps prevent models from rejecting text indiscriminately. We build on this insight: the goal is conditional text use, not blanket text removal. Our contribution is to replace broad SFT augmentation with model-specific preference pairs mined from actual text-following failures. Other recent conflict-oriented evaluations similarly study how VLMs behave when image and text cues disagree. Mixed Signals [10] constructs mismatched image-text pairs across domains and studies model biases under conflicting signals. These works motivate our focus on conflict resolution rather than ordinary VQA accuracy alone. In contrast to benchmark-only studies, we explicitly turn conflict errors into training pairs.

Document and visual question answering. DocVQA [7] evaluates question answering over document images, requiring OCR, layout understanding, and reasoning over visual text. VQAv2 [4] evaluates open-ended visual question answering with reduced language priors compared to earlier VQA datasets. We use DocVQA as the main testbed because conflict between visual document evidence and auxiliary text is central to the task. VQAv2 is used as a transfer check to test whether the mitigation is specific to document images. Text-centric VQA datasets such as TextVQA [13], ST-VQA [2], OCR-VQA [9], and InfographicVQA [8] emphasize reading and reasoning over text inside images. These tasks are related but not identical to our corrupted-auxiliary-text setting. In TextVQA, scene text is often the correct visual evidence; in our DocVQA conflict setting, auxiliary text can be an external shortcut that contradicts the image. This distinction explains why we report TextVQA as a boundary check rather than as the main benchmark.

Preference optimization. Preference optimization has become a common alternative to supervised fine-tuning for aligning generative models. DPO [11] directly optimizes a policy from chosen/rejected pairs without training an explicit reward model or running online RL. This makes it attractive for our setting because corrupted-text evaluation yields natural preference pairs. GRPO [12] is an online reinforcement learning method that estimates a group-relative baseline from multiple sampled responses. We compare DPO with GRPO-style optimization and find that free-form sampled-answer rewards are less effective for this failure

mode. Several recent works adapt preference optimization to multimodal models. mDPO [14] argues that standard multimodal DPO can suffer from unconditional preference learning, where the image condition is underused. Re-Align [16] constructs dual preference data with retrieval to incorporate visual preference signals. Task Preference Optimization [17] uses task-specific visual supervision to improve multimodal perception, and on-policy DPO has been explored for hallucination mitigation [18]. Our work is complementary: rather than building generic human or retrieval preferences, we mine preference pairs from a specific modality-conflict failure mode.

Modern VLMs. Qwen2-VL [15] and Qwen2.5-VL [1] emphasize dynamic-resolution visual processing and strong document/OCR capability. LLaVA-style models [5, 6] provide a widely used visual instruction tuning framework. Our transfer experiments suggest that these model families differ substantially in the base competence needed for text-bias correction.

3. Method

3.1. Problem Setup

Each example consists of an image x , a question q , auxiliary text t , and a ground-truth answer y . We consider three main textual conditions. In the *match* condition, the auxiliary text is consistent with the image. In the *irrelevant* condition, the text is unrelated. In the *corrupted* condition, the text contains a plausible but wrong answer. The desired behavior is not to ignore text unconditionally. Instead, the model should use text when it is reliable and fall back to visual evidence when text conflicts with the image.

3.2. Mining Text-Following Errors

We first evaluate a base VLM on corrupted examples. Let $\hat{y}$ be the model prediction. We define a *text-following error* when all of the following hold:

1. the model prediction is incorrect;
 2. the prediction is non-empty and does not match the ground truth;
 3. the prediction appears in the corrupted auxiliary text.
- For the third condition, we use normalized string matching. For multi-token answers, we also accept high token overlap with the auxiliary text. This filtering is important: an arbitrary wrong answer is not necessarily a text-bias error, but an incorrect prediction copied from corrupted text is a much stronger signal.

For each mined error, we construct a DPO pair:

$$y^+ = y, \quad y^- = \hat{y}, \quad (1)$$

where y^+ is the image-grounded answer and y^- is the model’s text-following wrong answer. We call this dataset

229 *text-following span DPO.*

230 **Why error mining matters.** The filtering step is not
231 merely a data-cleaning heuristic. It determines what behav-
232 ioral claim the training data supports. If we trained on every
233 wrong prediction, the rejected answer could correspond to
234 many different failure modes: OCR failure, answer-format
235 mismatch, arithmetic error, hallucination, or missing docu-
236 ment evidence. Such pairs would still be valid supervised
237 corrections, but they would not specifically target blind faith
238 in text. By requiring the wrong answer to appear in the cor-
239 rupted auxiliary text, we isolate a narrower intervention: the
240 model saw a misleading textual cue and produced that cue
241 as its answer. This makes the preference pair interpretable
242 as a conflict-resolution example rather than as generic answer
243 correction.

244 This distinction is important for our comparison with
245 SFT and GRPO: SFT teaches the correct answer, GRPO
246 must discover useful alternatives through sampling, while
247 our mined DPO pair directly asks the model to reverse a
248 concrete preference it already exhibited under conflict.

249 3.3. Why This Is Not Text Removal

250 A natural baseline is to delete the auxiliary text or instruct
251 the model to ignore it. However, this changes the task rather
252 than solving the conflict. In real document and visual-
253 question-answering settings, auxiliary text may come from
254 OCR, retrieved context, metadata, or captions. Some of it
255 is useful, some is irrelevant, and some may be wrong. The
256 model must therefore decide whether the text is supported
257 by the image, not merely whether text is present.

258 Our training data encodes this conditional behavior
259 through the surrounding conditions: corrupted examples
260 provide the main negative signal, match examples test reli-
261 able text use, and irrelevant/no-text examples test whether
262 the model can answer without overusing unsupported con-
263 text. Thus the desired policy is not

264 always distrust text, (2)

265 but rather

266 trust text only when it is visually grounded. (3)

267 This is why preservation metrics are central in our evalua-
268 tion. An approach that improves corrupted accuracy by dis-
269 carding all auxiliary text would be a weaker solution than
270 one that improves corrupted robustness while preserving
271 match and irrelevant performance.

272 3.4. DPO Objective

273 Given a policy π_θ and a reference policy π_{ref} , DPO opti-
274 mizes the preference that y^+ should be more likely than y^-

under the prompt (x, q, t) :

$$\mathcal{L}_{\text{DPO}} = -\log \sigma \left(\beta \left[\log \pi_\theta(y^+ | x, q, t) - \log \pi_{\text{ref}}(y^+ | x, q, t) \right. \right. \\ \left. \left. - \log \pi_\theta(y^- | x, q, t) + \log \pi_{\text{ref}}(y^- | x, q, t) \right] \right). \quad (4)$$

In our setting, this objective directly encodes the desired
correction: the image-grounded answer should be preferred
over the corrupted-text answer.

511 3.5. Multi-Condition Error-Mined DPO

Text-following span DPO uses only corrupted examples.
To reduce drift and improve stability, we also test a multi-
condition extension. We mine errors from corrupted, match,
irrelevant, and no-text evaluations. For corrupted examples,
we keep the strict text-following filter: the wrong predic-
tion must appear in corrupted auxiliary text. For match, ir-
relevant, and no-text examples, we use a broader condition-
specific error filter: the prediction must be incorrect, non-
empty, and distinct from the ground truth.

We assign lower weights to non-corrupted errors:

$$w_{\text{corrupted}} = 1.0, \quad w_{\text{match}} = 0.25, \\ w_{\text{irrelevant}} = 0.5, \quad w_{\text{no-text}} = 0.5. \quad (5)$$

This preserves the central text-bias signal while adding
weak regularization from other conditions.

512 3.6. GRPO-Style Baselines

We also test GRPO-style optimization. For each prompt,
the model samples multiple candidate answers. A reward
function gives positive reward for matching the chosen an-
swer and penalties for matching the rejected answer, mis-
leading spans mined from corrupted text, or any answer
copied from the auxiliary text. The reward also mildly en-
courages short answer-like completions. This tests whether
online sampled-answer optimization can discover the same
correction signal without explicit pairwise DPO.

513 4. Experiments

514 4.1. Metrics

We report three complementary metrics. *Corrupted accu-
racy* measures task performance when auxiliary text is mis-
leading. *Incorrect aux-hit* measures the fraction of incor-
rect predictions that appear in the corrupted auxiliary text.
Lower incorrect aux-hit indicates less text-following among
remaining errors, but it is not sufficient by itself: a model
could reduce aux-hit by hallucinating unrelated answers.
Therefore, we interpret it together with corrupted accuracy
and preservation on match/irrelevant conditions. We also
report corrected/regressed counts relative to the base model
when available.

4.2. Implementation Details

Unless otherwise stated, we train Qwen3-VL-8B-Instruct with LoRA adapters. The main DocVQA split uses 800 corrupted examples for mining and a disjoint held-out corrupted set of 200 examples for evaluation. The text-following span dataset is built only from base-model failures satisfying the corrupted-text overlap filter. For the initial span-DPO run, we use LoRA rank 8, LoRA alpha 16, learning rate 10^{-5} , DPO $\beta = 0.1$, and 700 DPO steps. The stronger rank-16 runs use LoRA alpha 32 with the same learning rate and number of steps. The multi-condition error-mined run uses 273 total preference rows, split into 245 train and 28 test rows, with source counts of 200 corrupted, 7 match, 33 irrelevant, and 33 no-text errors. It trains for 900 DPO steps with LoRA rank 16, alpha 32, learning rate 10^{-5} , and $\beta = 0.1$.

All reported transfer experiments use the same evaluation code and the same model family-specific prompt template used during mining. This matters because our LLaVA-Next experiments show strong sensitivity to task and template compatibility. We therefore interpret LLaVA results as boundary evidence rather than as a claim that the method is model-agnostic.

Corruption construction. For DocVQA, corrupted auxiliary text is created by replacing the answer-bearing textual hint with a plausible but incorrect answer span. The corrupted span is intentionally answer-like rather than random noise. This makes the setting difficult: a model that relies on surface text can produce a fluent and plausible answer, while a model that checks the image must reject the hint. For match and irrelevant conditions, the auxiliary text is either consistent with the ground-truth answer or unrelated to the queried evidence. Together, these conditions separate three behaviors that would be conflated in ordinary VQA accuracy: using helpful text, ignoring irrelevant text, and correcting misleading text.

We use exact normalized answer matching for VQA accuracy. For incorrect aux-hit, we measure whether an incorrect prediction overlaps the corrupted auxiliary answer span after normalization. Because aux-hit is computed only among wrong predictions, it diagnoses residual failure type rather than replacing accuracy.

4.3. DocVQA Main Results

DocVQA is our main testbed. We use Qwen3-VL-8B-Instruct as the primary model and evaluate on a held-out corrupted split. Table 1 shows that text-following span DPO substantially improves both accuracy and text-bias behavior. Increasing LoRA rank from 8 to 16 improves the result further, while the multi-condition error-mined variant gives the best overall performance.

Table 1. DocVQA corrupted held-out results with Qwen3-VL-8B.

Method	Acc.	Incorrect aux-hit	Corr./Reg.
Base	0.715	89.5%	—
TF-span DPO, r8	0.900	55.0%	37 / 0
TF-span DPO, r16	0.915	47.1%	40 / 0
Multi-cond. error DPO	0.920	43.8%	41 / 0

Table 2. DocVQA DPO hyperparameter grid. Higher LoRA capacity matters more than β in this setting.

β	Rank	Acc.	Aux-hit	Corr./Reg.	Match/Irr.
0.1	8	0.900	55.0%	37 / 0	1.00 / 0.96
0.2	8	0.900	55.0%	37 / 0	0.98 / 0.91
0.1	16	0.915	47.1%	40 / 0	0.98 / 0.91
0.2	16	0.915	47.1%	40 / 0	0.98 / 0.91

The multi-condition result suggests that the highest-performing variant is not corrupted-only. However, the corrupted filter remains essential: corrupted rows still require the prediction to be copied from corrupted auxiliary text. The additional match, irrelevant, and no-text errors act as weak condition-specific regularizers rather than replacing the text-following signal.

The main DocVQA result is encouraging because it increases corrupted accuracy and lowers aux-hit without large match/irrelevant degradation.

4.4. Hyperparameter and Construction Ablations

We run a small 2×2 grid over DPO β and LoRA rank on the same DocVQA split. Table 2 shows that increasing LoRA rank from 8 to 16 is more useful than increasing β from 0.1 to 0.2. Rank 16 improves corrupted accuracy from 0.900 to 0.915 and reduces incorrect aux-hit from 55.0% to 47.1%, while changing β does not change the corrupted metric in this grid.

Prompt-augmented counterfactual DPO, which reuses the same rejected answer across corrupted, no-text, match, and irrelevant wrappers, ties plain rank-16 DPO at 0.915 accuracy and 47.1% aux-hit. Multi-condition error-mined DPO instead uses the actual failed prediction from each condition as the rejected answer, improving to 0.920 accuracy and 43.8% aux-hit.

4.5. DPO versus GRPO

We compare DPO with GRPO-only and DPO-to-GRPO variants on the same disjoint DocVQA split. GRPO samples multiple answers per prompt and optimizes a reward that encourages correctness and penalizes copying corrupted text. Despite this denser reward, GRPO does not move the held-out greedy behavior meaningfully. Table 3 shows that DPO is the effective intervention in our setting.

Table 3. GRPO-style optimization does not improve over DPO on DocVQA.

Variant	Init	Acc.	Incorrect aux-hit
Base	–	0.715	89.5%
GRPO-only, 100 steps	base	0.710	89.7%
GRPO-only, 300 steps	base	0.705	89.8%
GRPO-only, dense reward	base	0.705	89.8%
TF-span DPO	base	0.900	55.0%
DPO → GRPO	DPO	0.900	60.0%

DPO-to-GRPO changes only a few held-out predictions relative to DPO and does not improve accuracy. This supports our hypothesis: when high-precision chosen/rejected pairs are available, direct pairwise optimization provides a clearer learning signal than free-form sampled-answer rewards.

4.6. SFT: Prior Mitigation and Our Ablation

The original Words or Vision study already evaluates SFT with text augmentation as a mitigation strategy [3]. That SFT setting uses a broader training mixture, including original VQA, text-only, match, corruption, and irrelevance samples, and shows that SFT can reduce text reliance while maintaining the ability to benefit from matched text. This is an important baseline and motivates our own mitigation study.

Our experiment is different. We do not reproduce the full augmented SFT recipe from Deng et al. [3]. Instead, we test a narrower question: after constructing text-following DPO pairs, can an auxiliary SFT preservation loss on match/irrelevant examples further stabilize the model? This is a preservation ablation, not a claim that SFT in general is ineffective. SFT is plausible because it directly teaches the model to output correct answers. However, the central error in our setting is not simply that the model does not know the answer; it is that the model prefers a misleading textual answer over visual evidence. SFT provides only a positive target, while DPO explicitly contrasts the image-grounded answer against the text-following wrong answer.

Our SFT experiment uses auxiliary preservation data from match and irrelevant conditions and mixes it with span DPO. As shown in Table 4, this does not improve over pure DPO on VQAv2: corrupted accuracy decreases from 0.645 to 0.635, incorrect aux-hit increases from 40.8% to 42.5%, and match/irrelevant preservation remains unchanged. Thus, in our experiments, auxiliary SFT preservation acts as a negative ablation, while the prior full SFT augmentation result remains a valid mitigation baseline from the original work.

Table 4. VQAv2 transfer results with Qwen3-VL-8B.

Variant	Corr. Acc.	Aux-hit	Match	Irr.
Base	0.545	65.9%	0.92	0.77
Span DPO	0.645	40.8%	0.90	0.75
Span DPO + SFT	0.635	42.5%	0.90	0.75

Table 5. TextVQA question-type matched answer-hint transfer. DocVQA DPO transfers weakly to Qwen models but not reliably to LLaVA-Next.

Model	Base Acc.	DPO Acc.	Base aux-hit	DPO aux-hit
Qwen3-VL-8B	0.821	0.827	11.5%	8.3%
Qwen2-VL-7B	0.832	0.845	23.1%	8.3%
LLaVA-N-7B	0.535	0.443	50.0%	62.0%
LLaVA-N-13B	0.612	0.602	35.7%	28.8%

4.7. VQAv2 Transfer Check

We repeat the disjoint split protocol on VQAv2. The effect is positive but weaker than on DocVQA. As shown in Table 4, span DPO improves corrupted accuracy by 10 points and reduces incorrect aux-hit, but match and irrelevant preservation drop slightly. Adding auxiliary SFT preservation does not recover this drop and slightly worsens corrupted performance.

4.8. TextVQA Boundary Check

We also test whether a DocVQA-trained span DPO adapter transfers to TextVQA. TextVQA is useful because it is text-centric, but it differs from DocVQA conflict: scene text is often the correct visual evidence rather than an external auxiliary shortcut. We therefore construct several TextVQA corruption variants, including same-image OCR distractors, cross-sample answers, and question-type matched answer hints.

Table 5 shows the strongest TextVQA conflict variant, question-type matched answer hints, on 200 examples. For Qwen3-VL, the base model already has low incorrect aux-hit under corrupted TextVQA hints, and DocVQA DPO gives only a small gain. Qwen2-VL shows a similar but slightly stronger positive direction. LLaVA-Next behaves differently: the conflict is stronger, but DocVQA DPO does not reliably fix it and can worsen corrupted accuracy.

These results sharpen the scope of our claim. The method is not a universal accuracy booster for all text-heavy VQA tasks. It is most effective when the base model actually exhibits text-following failures under corrupted auxiliary text and already has enough visual/OCR competence to answer from the image.

4.9. Model Transfer

We test whether the same mitigation extends beyond Qwen3-VL. Table 6 summarizes Qwen2-VL-7B and

SFT and DPO reduce blind faith in text through different supervision signals




Figure 2. Comparison of mitigation signals and representative outcomes. The SFT panel reports the prior Words or Vision text-augmentation result on VQAv2 with Qwen2-VL-7B, while the DPO panel reports our DocVQA results with Qwen3-VL-8B. The figure is therefore not a direct dataset-matched comparison; it illustrates that SFT can mitigate text bias through broad supervised augmentation, whereas our DPO approach targets model-specific text-following errors through pairwise preference learning.

Table 6. Model transfer on DocVQA. DPO helps Qwen2-VL but not LLaVA-Next under our setting.

Model	Setting	Acc.	Aux-hit	Corr./Reg.
Qwen2-VL-7B	base	0.570	93.0%	–
Qwen2-VL-7B	span DPO	0.775	86.7%	41 / 0
LLaVA-N-7B	base	0.105	90.5%	–
LLaVA-N-7B	span DPO	0.100	91.7%	1 / 2
LLaVA-N-13B	base	0.115	92.7%	–
LLaVA-N-13B	span DPO	0.130	95.4%	3 / 0

LLaVA-Next results on DocVQA. Qwen2-VL benefits substantially, improving from 0.570 to 0.775 corrupted accuracy. LLaVA-Next, however, has very low base DocVQA accuracy under our prompt/template, and DPO provides little benefit.

These results show that the method is not simply a universal post-hoc fix. It requires a base model that can already solve the task visually; DPO then teaches the model to resist misleading text.

5. Analysis and Discussion

5.1. Why DPO Fits This Problem

The corrupted-text setting gives unusually clean preference pairs. For many failures, the model’s wrong answer is directly copied from the corrupted auxiliary text. Thus the comparison is explicit: the image-grounded answer should be preferred over the text-following wrong answer. DPO optimizes exactly this comparison. This complements, rather than invalidates, the SFT mitigation in Words or Vi-

sion [3]. SFT with text augmentation teaches the model from a balanced mixture of cases and can reduce broad text reliance. Our DPO variant instead uses a narrower but sharper signal: for each mined failure, it tells the model which specific text-derived answer should lose to the visual answer. In contrast, GRPO must first sample useful candidate answers and then rely on a hand-designed reward. If the sampled candidates rarely contain the correct answer, or if most candidates are similar text-following errors, the relative reward signal is weak. This explains why GRPO logs can show diagnostic reward events while greedy held-out predictions barely change.

Our DPO construction uses the benchmark answer as the chosen response. This is appropriate because the task is not open-ended preference alignment but targeted VQA correction: the correct answer should beat an incorrect answer copied from misleading text. The advantage over SFT is contrastive. SFT says only what to output, while DPO also identifies the tempting text-derived answer that should be suppressed.

5.2. Why Not Simply Ignore the Text?

Ignoring auxiliary text is an appealing but incomplete baseline. It would avoid some corrupted-text failures, but it does not solve the actual decision problem faced by a VLM. At inference time, the model is not given a label saying whether the text is matched, irrelevant, or corrupted. If the model learns to discard text wholesale, it can lose useful information in match cases and may fail to exploit reliable OCR or retrieved context.

Our objective is therefore not “use less text” but “use text conditionally”. The model should keep using textual con-

text when it is visually supported and reject it when it conflicts with the image. DPO is appropriate because it changes the model’s preference in exactly the conflict cases we can identify:

image-grounded answer $\succ$ text-following wrong answer. (6)

This is different from input ablation or prompt-level instruction to ignore text. The latter can reduce exposure to misleading text, but it does not teach the model which answer should be preferred when both visual evidence and misleading text are present. Our use of match and irrelevant preservation metrics is intended to check precisely this point: a successful method should improve corrupted robustness without collapsing into text avoidance.

5.3. What Counts as Solving Text Bias?

We therefore define a successful correction operationally through four criteria: held-out corrupted accuracy should improve, incorrect aux-hit should decrease, match/irrelevant settings should be preserved, and the trend should be at least partially reproducible across nearby settings such as DPO hyperparameters or related VQA benchmarks. Our DocVQA results satisfy the first three criteria strongly and the fourth criterion partially: Qwen-family transfer is positive, while LLaVA and TextVQA reveal clear boundary conditions.

5.4. Why Multi-Condition Error Mining Helps

Text-following span DPO is the core filtering strategy, but the best result comes from multi-condition error-mined DPO. This does not mean that unfiltered errors are sufficient. Rather, corrupted errors remain strictly filtered for text-following behavior, while other conditions contribute lower-weight stability signals. This helps prevent the model from over-specializing to the corrupted prompt format. The gain is modest because non-corrupted failures are sparse in DocVQA, but it consistently improves the main held-out metric.

5.5. Boundary Conditions

Our transfer experiments reveal a clear boundary condition. The method works when the base VLM already has sufficient visual and OCR competence. Qwen-family models satisfy this condition on DocVQA and improve substantially. LLaVA-Next models, under our prompt/template, have very low base DocVQA accuracy. In this regime, DPO cannot create document-reading ability from scratch; it can only alter preferences among answers the model can already represent.

TextVQA experiments lead to a related conclusion. DocVQA-trained DPO transfers weakly to TextVQA and does not reliably fix LLaVA-specific TextVQA conflict failures. Even after mining larger LLaVA-specific TextVQA

DPO data and adding strict+preserve pairs, corrupted performance does not improve. This suggests that TextVQA conflict involves different OCR and answer-format dynamics than DocVQA. Therefore, we treat TextVQA and LLaVA results as boundary analyses rather than main success claims.

5.6. Limitations

Our study has several limitations. First, the text-following filter is operational rather than causal; accidental overlap with corrupted text is possible. Second, most experiments use one held-out split and limited seeds due to compute constraints. Third, our GRPO implementation uses free-form sampled answers, and candidate-constrained variants may be stronger. Finally, the method assumes ground-truth answers for preference construction, making it suitable for supervised mitigation but not directly for fully unlabeled deployment.

6. Conclusion

We studied preference optimization as a mitigation for VLM text bias under corrupted auxiliary text. By mining text-following errors and training with DPO, we convert a diagnostic failure mode into a direct preference-learning signal. On DocVQA, DPO improves corrupted-text accuracy and reduces copying from misleading text.

References

- [1] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, et al. Qwen2.5-vl technical report. *arXiv preprint arXiv:2502.13923*, 2025. 3
- [2] Ali Furkan Biten, Ruben Tito, Andres Mafla, Lluís Gomez, Marçal Rusiñol, Ernest Valveny, C. V. Jawahar, and Dimosthenis Karatzas. Scene text visual question answering. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2019. 3
- [3] Ailin Deng, Tri Cao, Zhirui Chen, and Bryan Hooi. Words or vision: Do vision-language models have blind faith in text? In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025. 1, 2, 3, 6, 7
- [4] Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2017. 3
- [5] Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. *Advances in Neural Information Processing Systems*, 2023. 3
- [6] Haotian Liu, Chunyuan Li, Yuheng Li, Bo Li, Yuanhan Zhang, Sheng Shen, and Yong Jae Lee. Llava-next: Improved reasoning, ocr, and world knowledge. *LLaVA project blog*, 2024. 3

- [7] Minesh Mathew, Dimosthenis Karatzas, and C. V. Jawahar. Docvqa: A dataset for vqa on document images. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2021. 3
- [8] Minesh Mathew, Viraj Bagal, Ruben Tito, Dimosthenis Karatzas, Ernest Valveny, and C. V. Jawahar. Infographicvqa. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2022. 3
- [9] Anand Mishra, Shashank Shekhar, Ajeet Kumar Singh, and Anirban Chakraborty. Ocr-vqa: Visual question answering by reading text in images. In *International Conference on Document Analysis and Recognition*, 2019. 3
- [10] Pouya Pezeshkpour, Moin Aminnaseri, and Estevam Hruschka. Mixed signals: Decoding vlms’ reasoning and underlying bias in vision-language conflict. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 24833–24848, 2025. 3
- [11] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, 2023. 1, 3
- [12] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. 1, 3
- [13] Amanpreet Singh, Vivek Natarajan, Meet Shah, Yu Jiang, Xinlei Chen, Dhruv Batra, Devi Parikh, and Marcus Rohrbach. Towards vqa models that can read. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019. 3
- [14] Fei Wang, Wenxuan Zhou, James Y. Huang, Nan Xu, Sheng Zhang, Hoifung Poon, and Muhao Chen. mdpo: Conditional preference optimization for multimodal large language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024. 3
- [15] Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, et al. Qwen2-vl: Enhancing vision-language model’s perception of the world at any resolution. *arXiv preprint arXiv:2409.12191*, 2024. 3
- [16] Shuo Xing, Yuping Wang, Peiran Li, Ruizheng Bai, Yueqi Wang, Chengxuan Qian, Huaxiu Yao, and Zhengzhong Tu. Re-align: Aligning vision language models via retrieval-augmented direct preference optimization. *arXiv preprint arXiv:2502.13146*, 2025. 3
- [17] Ziang Yan, Zhilin Li, Yinan He, Chenting Wang, Kunchang Li, Xinhao Li, Xiangyu Zeng, Zilei Wang, Yali Wang, Yu Qiao, Limin Wang, and Yi Wang. Task preference optimization: Improving multimodal large language models with vision task alignment. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025. 3
- [18] Zhihe Yang, Xufang Luo, Dongqi Han, Yunjian Xu, and Dongsheng Li. Mitigating hallucinations in large vision-language models via dpo: On-policy data hold the key. In

Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2025. 3

677
678